

CS 584 Natural Language Processing

Yue Ning
Department of Computer Science
Stevens Institute of Technology

Lecture 3: Word vectors

Learning Objectives

- ❖ Understand word2vec (Skip-gram and Glove models).
- ❖ Implement gradient descent for skip-gram
- ❖ Understand evaluation for word vectors
- ❖ Apply word2vec to analyze text

What⁺_o is the meaning of
a word?




Represent the meaning of a word

- ❖ Definition: **meaning** (webster dictionary)
 - The idea that is represented by a word, phrase, etc
 - The idea that a person wants to express by using words, signs, etc
 - The idea that is expressed in a word of writing, art, etc.

What does it mean to “represent meaning”?

Represent the meaning of a word

Knowledge-based representation



One-hot representation


Issues: difficult to compute the similarity (i.e. comparing "car" and "motorcycle")

$$\begin{matrix} [0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & \dots & 0] & \text{AND} & [0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & \dots & 0] \\ \text{car} & & & & & & & & & & \text{motorcycle} \end{matrix} = 0$$

Co-occurrence Matrix

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Low-dimensional dense word vector (e.g., by dimension reduction on the co-occurrence matrix)





Why discrete representations fail

(WordNet, one-hot, TF-IDF)



Representing words with Knowledge-base

- ❖ Common solution: use knowledge base, e.g. WordNet, a thesaurus containing lists of synonym sets and hypernyms (“is a” relationships)

e.g. synonym sets containing “good”:

```
from nltk.corpus import wordnet as wn
poses = { 'n': 'noun', 'v': 'verb', 's': 'adj (s)', 'a': 'adj', 'r': 'adv' }
for synset in wn.synsets("good"):
    print("{}: {}".format(poses[synset.pos()],
        ", ".join([l.name() for l in synset.lemmas()])))
```

```
noun: good
noun: good, goodness
noun: good, goodness
noun: commodity, trade_good, good
adj: good
adj (sat): full, good
adj: good
adj (sat): estimable, good, honorable, respectable
adj (sat): beneficial, good
adj (sat): good
adj (sat): good, just, upright
...
adverb: well, good
adverb: thoroughly, soundly, good
```

e.g. hypernyms of “panda”:

```
from nltk.corpus import wordnet as wn
panda = wn.synset("panda.n.01")
hyper = lambda s: s.hypernyms()
list(panda.closure(hyper))
```

```
[Synset('procyonid.n.01'),
Synset('carnivore.n.01'),
Synset('placental.n.01'),
Synset('mammal.n.01'),
Synset('vertebrate.n.01'),
Synset('chordate.n.01'),
Synset('animal.n.01'),
Synset('organism.n.01'),
Synset('living_thing.n.01'),
Synset('whole.n.02'),
Synset('object.n.01'),
Synset('physical_entity.n.01'),
Synset('entity.n.01')]
```

Limitations with WordNet

- ❖ Great as a resource but missing nuance
 - e.g. “proficient” is listed as a synonym for “good”. This is only correct in some contexts.
- ❖ Missing new meanings of words
 - e.g., wicked, badass, nifty, wizard, genius, ninja, bombest
Impossible to keep up-to-date!
- ❖ Subjective
- ❖ Requires human labor to create and adapt
- ❖ Can’t compute accurate word similarity

Similarity Metrics of Vectors

❖ Cosine similarity

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

❖ Inner product (dot product)

Representing words as one-hot vectors

- ❖ In traditional NLP, we regard words as **discrete symbols**: hotels, conference, motel - a local list representation

Means one 1, the
rest 0s

- ❖ Words can be represented by **one-hot** vectors:
 - motel = [0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
 - hotel = [0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0]
- ❖ Vector dimension = number of words in vocabulary (e.g. 500, 000)
- ❖ Limitations: High dimensional, sparsity leading to inefficiency in storage and computation, lack of contextual information, no quantification of similarity, cannot handle out-of-vocabulary or unseen words

Representing documents as bag of words

- ❖ Bag of words: count, TFIDF, etc
- ❖ Documents can be represented by vectors:
 - $\text{doc1} = [0\ 0\ 3\ 0\ 2\ 0\ 0\ 0\ 0\ 0\ 5\ 0\ 0\ 1\ 0]$
 - $\text{doc2} = [0\ 1\ 0\ 0\ 0\ 2\ 0\ 0\ 4\ 0\ 0\ 1\ 1\ 0\ 0]$
- ❖ Vector dimension = number of words in vocabulary (e.g. 500, 000)

TF-IDF

- ❑ Term frequency - inverse document frequency
 - ❑ First define a **vocabulary** with length V
 - ❑ For **each word** in the vocabulary, calculate the vector for **current document**:
 $[TF_1 / DF_1, \dots, TF_v / DF_v, \dots, TF_V / DF_V]$
 - ❑ Where, **TF** is the frequency of word v appearing in the current document, and
 - ❑ **DF** is the # of documents where word v appears.
 - ❑ Sometimes, you need to add **smoothing** when **DF** is a small value

Limitations with words as discrete symbols

❖ Example: in web search, if user searches for “Seattle motel”, we would like to match documents containing “Seattle hotel”

❖ But

- motel = [0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]

- hotel = [0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0]

These two vectors are **orthogonal**. There is no natural notion of **similarity** of one-hot vectors.

❖ Solution:

- Could try to rely on WordNet’s list of synonyms to get similarity?
 - but it is well-known to fail badly: incompleteness, etc
 - instead: **learn to encode similarity in the vectors themselves**



Count-based distributional semantics

(co-occurrence, SVD)



Representing with co-occurrence matrix

With a co-occurrence matrix X

- 2 options: windows vs. full document
- **Window**: use window around each word -> captures both syntactic (POS) and semantic information (similar to word2vec)
- **Word-document** co-occurrence matrix: give general topics (all sports terms will have similar entries), leading to “latent semantic analysis”

Window-based co-occurrence matrix

Window length 1 (more common: 5-10)

- **Symmetric** (irrelevant whether left or right context)
- Example corpus:
 - I like deep learning.
 - I like NLP.
 - I enjoy flying.

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Limitations with simple co-occurrence vectors

Limitations:


- Increase in size with vocabulary
- Very high dimensional: requires a lot of storage
- Subsequent classification models have sparsity issues
- Models are less robust

Solution: low dimensional dense vectors

- Idea: store “most” of the important information in a fixed, small number of dimensions: **a dense vector**
- Usually 25-1000 dimensions
- How to reduce the dimensionality?

Method 1: dimensionality reduction on X

- Singular value decomposition (SVD) of co-occurrence matrix X
- Factorizes X into $U\Sigma V^T$, where U and V are orthonormal



- Retain only k singular values, in order to generalize.
- A is the best rank k approximation to X, in terms of least squares.
- Classic linear algebra result.
- Expensive to compute for large matrices.

Simple SVD word vectors in Python


- corpus:
 - I like deep learning. I like NLP. I enjoy flying.

```
import numpy as np
la = np.linalg
words = ["I", "like", "enjoy",
         "deep", "learnig", "NLP", "flying", "."]
X = np.array([[0,2,1,0,0,0,0,0],
              [2,0,0,1,0,1,0,0],
              [1,0,0,0,0,0,1,0],
              [0,1,0,0,1,0,0,0],
              [0,0,0,1,0,0,0,1],
              [0,1,0,0,0,0,0,1],
              [0,0,1,0,0,0,0,1],
              [0,0,0,0,1,1,1,0]])

U, s, Vh = la.svd(X, full_matrices=False)
```

Simple SVD word vectors in Python

- Printing first two columns of U corresponding to the 2 biggest singular values.



Running an SVD on raw counts doesn't work well.

Hacks to X

(several used in [Rohde et al. 2005](#))

- **Scaling the counts** in the cells can help a lot.
- Problem: function words (the, he, has) are too frequent, thus syntax has too much impact. Some fixes:
 - $\min(X, t)$, with $t \approx 100$
 - Ignore the function words
- Ramped windows that count closer words more (weighted sum of all occurrences of b in proximity to a)
- Use Pearson correlations instead of counts, then set negative values to 0

Interesting syntactic patterns emerge in the vectors



from “*An Improved Model of Semantic Similarity Based on Lexical Co-Occurrence*”. Rohde et al., 2005

Limitations of SVD

- Computational cost scales quadratically
- Bad for millions of words or documents
- Hard to incorporate new words or documents
- Different learning regime than other DL models


⁺ ○ Prediction-based embeddings

(Skip-gram, CBOW, Glove)



Representing words by their **CONTEXT**

- ❖ **Distributional semantics**: A word's meaning is given by the words that frequently appear close-by.
- ❖ One of the most successful ideas of modern statistical NLP!
- ❖ When a word w appears in a text, its **context** is the set of words that appear nearby (within a fixed-size window).
- ❖ Use the many contexts of w to build up a representation of w



Word vectors


- ❖ We will build a dense vector for each word, chosen so that the vectors of words that appear in similar context will be similar.

$$\text{banking} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

- ❖ Note: word vectors are sometimes called word embeddings or word representations. They are a distributed representation.

Word meaning as a word vector

- visualization




<https://swatimeena989.medium.com/training-word2vec-using-gensim-14433890e8e4>

Word2vec: Overview

- **Word2vec** (Mikolov et al. 2013) is a framework for learning word vectors
- Idea:
 - We have a large corpus of text
 - Every word in a fixed vocabulary is represented by a vector
 - Go through each position t in the text, which has a center word c and context (“outside”) words o
 - Use the **similarity of the word vectors** for c and o to **calculate the probability** of o given c (or vice versa)
 - Keep **adjusting** the word vectors to maximize this probability


Word2vec: Overview

- Example windows and process for computing $P(w_{t+j}|w_t)$



Word2vec: Overview

- Example windows and process for computing $P(w_{t+j}|w_t)$



Word2vec: objective function

- For each position $t = 1, \dots, T$, predict context words within a window of fixed size m , given center word w_t .

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{-m \leq j \leq m, j \neq 0} P(w_{t+j} | w_t; \theta)$$

- The **objective function** J is the (average) negative log likelihood

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} | w_t; \theta)$$

Minimizing objective function $\Leftrightarrow$ Maximizing predictive accuracy

At this point, what are we actually learning?

Word2vec overview with vectors

- We want to minimize the objective function:


$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} | w_t; \theta)$$

- Question: how to calculate $P(w_{t+j} | w_t; \theta)$?
- Answer: we will **use two vectors per word w** :
 - v_w when w is a **center** word
 - u_w when w is a **context** word
- Then for a center word c and a context word o :

$$P(o|c) = \frac{\exp(u_o^\top v_c)}{\sum_{w \in V} \exp(u_w^\top v_c)}$$

Word2vec: objective function

- Example windows and process for computing $P(w_{t+j}|w_t)$
- $P(u_{\text{problems}} | v_{\text{into}})$ short for $P(\text{problems}|\text{into}; u_{\text{problems}}, v_{\text{into}}, \theta)$



Word2vec: prediction function

Exponentiation makes anything positive

Dot product compares similarity of o and c. larger dot product = larger probability

$$P(o|c) = \frac{\exp(u_o^\top v_c)}{\sum_{w \in V} \exp(u_w^\top v_c)}$$

Normalize over entire vocabulary to give probability distribution


- This is an example of the softmax function

$$\text{softmax}(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)} = p_i$$

- The softmax function maps arbitrary values x_i to a probability distribution p_i
 - “max” because amplifies probability of largest x_i
 - “soft” because still assigns some probability to smaller x_i
 - frequently used in deep learning

Training a model by optimizing parameters

- To train a model, we adjust parameters to minimize a loss e.g. below, for a simple convex function over two parameters. Contour lines show levels of objective function



To train the model: compute all vector gradients

- Recall: θ represents all model parameters, in one long vector
- In our case, with d -dimensional vectors and V -many words:

$$\theta = \begin{bmatrix} v_{aardvark} \\ v_a \\ \vdots \\ v_{zebra} \\ u_{aardvark} \\ u_a \\ \vdots \\ u_{zebra} \end{bmatrix} \in \mathbb{R}^{2dV}$$

- Remember: every word has two vectors
- We optimize these parameters by walking down the gradient

Word2vec derivations of gradient

- Useful basics:

$$\frac{\partial \mathbf{x}^\top \mathbf{a}}{\partial \mathbf{x}} = \frac{\partial \mathbf{a}^\top \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$$

- If in doubt: write out with indices
- Chain rule: if $y = f(u)$ and $u = g(x)$, i.e. $y = f(g(x))$, then

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$

Exercise

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} | w_t; \theta)$$

Let's derive gradient for center word


Note: 1) For one example in window and one example outside window:

$$\log P(o|c) = \log \frac{\exp(u_o^\top v_c)}{\sum_{w=1}^V \exp(u_w^\top v_c)}$$

2) You then also need the gradient for context words (it's similar). That's all of the parameter θ here.

Calculating all gradients!

- We went through gradient for each **center vector** v in a window
- We also need gradients for **outside vectors** u
- Generally, in each window we will compute updates for all parameters that are being used in that window. For example:



Word2vec: more variations

- Why two vectors?
 - > easier optimization, average both at the end
- Two model variants:
 - **Skip-grams (SG)**: predict context (“outside”) words (position independent) given center word
 - **Continuous Bag of Words (CBOW)**: predict center word from context words
 - This lecture so far: **skip-gram model (SG)**

Stochastic Gradient Descent

- Problem: $J(\theta)$ is a function of all windows in the corpus (potentially billions!)
 - so $\nabla_{\theta}J(\theta)$ is very expensive to compute
- You would wait a long time before making a single update
- Very bad idea for pretty much all neural nets!
- Solution: Stochastic gradient descent (SGD)
 - Repeatedly sample windows and update after each sample
- Algorithm:

```
while True:
    window = sample_window(corpus)
    theta_grad = evaluate_gradient(J, window, theta)
    theta = theta - alpha * theta_grad
```

Stochastic gradients with word vectors

- Iteratively take gradients at each such window for SGD
- But in each window, we only have at most $2m+1$ words, so $\nabla_{\theta} J(\theta)$ is sparse!

$$\nabla_{\theta} J_t(\theta) = \begin{bmatrix} 0 \\ \vdots \\ \nabla_{v_{like}} \\ \vdots \\ 0 \\ \nabla_{u_I} \\ \vdots \\ \nabla_{u_{learning}} \\ \vdots \end{bmatrix} \in \mathbb{R}^{2dV}$$

Skip-gram model with negative sampling

- The normalization factor is too computationally expensive.

$$P(o|c) = \frac{\exp(u_o^\top v_c)}{\sum_{w \in V} \exp(u_w^\top v_c)}$$

- Hence, in standard word2vec you implement the skip-gram model with **negative sampling**
- Main idea: train binary logistic regressions for a **true/positive pair** (center word and word in its context window) versus several **noise/negative pairs** (the center word paired with a random word)


Skip-gram model with negative sampling

- From paper: “Distributed Representations of Words and Phrases and their Compositionality” (Mikolov et al. 2013)
- Overall objective function (maximize):

$$J_t(\theta) = \log \sigma(u_o^\top v_c) + \sum_{j=1}^k \mathbb{E}_{j \sim P(w)} [\log \sigma(-u_j^\top v_c)] \quad J(\theta) = \frac{1}{T} \sum_{t=1}^T J_t(\theta)$$

- The sigmoid function:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



- We maximize the probability of two words co-occurring in first log

Skip-gram model with negative sampling

- We take k negative samples (using word probabilities)
- Maximize probability that real outside word appears, minimize probability that random words appear around center word
- Noise distribution: $P(w) = U(w)^{3/4} / Z$, the unigram distribution $U(w)$ raised to the $3/4$ power outperformed other distributions significantly
- The power makes less frequent words be sampled more often

Count based vs. direct prediction based word vectors

- LSA, HAL (Lund & Burgess),
- COALS, Hellinger-PCA (Rohde et al, Lebrete & Collobert)

- Fast training
- Efficient usage of statistics
- Primarily used to capture word similarity
- Disproportionate importance given to large counts

- Skip-gram/CBOW (Mikolov et al)
- NNLM, HLBL, RNN (Bengio et al; Collobert & Weston; Huang et al; Mnih & Hinton)

- Scales with corpus size
- Inefficient usage of statistics
- Generate improved performance on other tasks
- Can capture complex patterns beyond word similarity

Encoding meaning in vector differences

[Pennington, Socher, and Manning, EMNLP 2014]

- Crucial insight: Ratios of co-occurrence probabilities can encode meaning components

	$x = \text{solid}$	$x = \text{gas}$	$x = \text{water}$	$x = \text{random}$
$P(x \text{ice})$	large	small	large	small
$P(x \text{steam})$	small	large	large	small
$\frac{P(x \text{ice})}{P(x \text{steam})}$	large	small	~ 1	~ 1

Encoding meaning in vector differences

[Pennington, Socher, and Manning, EMNLP 2014]

- Crucial insight: Ratios of co-occurrence probabilities can encode meaning components

	$x = \text{solid}$	$x = \text{gas}$	$x = \text{water}$	$x = \text{fashion}$
$P(x \text{ice})$	1.9×10^{-4}	6.6×10^{-5}	3.0×10^{-3}	1.7×10^{-5}
$P(x \text{steam})$	2.2×10^{-5}	7.8×10^{-4}	2.2×10^{-3}	1.8×10^{-5}
$\frac{P(x \text{ice})}{P(x \text{steam})}$	8.9	8.5×10^{-2}	1.36	0.96

Encoding meaning in vector differences

[Pennington, Socher, and Manning, EMNLP 2014]

- Q: How can we capture ratios of co-occurrence probabilities as linear meaning components in a word vector space?
- A: log-bilinear model with vector differences

$$w_i \cdot w_j = \log P(j|i)$$


$$(w_a - w_b) \cdot w_x = \log \frac{P(x|a)}{P(x|b)}$$

Combining the best of both worlds: Glove [Pennington et al., EMNLP 2014]

$$w_i \cdot w_j = \log P(j|i)$$

$$J = \sum_{i,j=1}^V f(X_{ij})(w_i^\top \hat{w}_j + b_i + \hat{b}_j - \log X_{ij})^2$$

- Fast training
- Scalable to huge corpora
- Good performance even with small corpus and small vectors



Glove results

[Pennington et al., EMNLP 2014]

Nearest words to frog:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



leptodactylidae



rana



eleutherodactylus



How to evaluate word vectors?



How to evaluate word vectors?

- Related to general evaluation in NLP: **Intrinsic** vs **extrinsic**
- **Intrinsic:**
 - Evaluation on a specific/intermediate subtask
 - Fast to compute
 - Helps to understand that system that generates the word vectors
 - Not clear if really helpful, unless correlation to real task is established
- **Extrinsic:**
 - Evaluation on a real task
 - Can take a long time to compute accuracy
 - Unclear if the sub system is the problem, or its interaction, or other subsystems
 - If replacing exactly one subsystem with another improves accuracy, Winning!


Intrinsic evaluation: Word vector analogy

- Word vector analogies

$$\boxed{a:b :: c:?} \longrightarrow \boxed{d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}}$$


man:woman :: king:?

- Evaluate word vectors by how well their **cosine distance** after addition captures intuitive semantic and syntactic analogy
- Identify the word vector which maximizes the cosine similarity
- Using intrinsic evaluation techniques should be handled with care




City 1 : State containing City 1 :: City 2 : State containing City 2

Glove visualizations: linear word analogy



Glove visualizations: Company - CEO



Intrinsic evaluation: Correlation evaluation

- Word vector distances and their correlation with human judgments
- Example dataset: WordSim353
 - <http://alfonseca.org/eng/research/wordsim353.html>

Word 1	Word 2	Human (mean)
tiger	cat	7.35
tiger	tiger	10
book	paper	7.46
computer	internet	7.58
plane	car	5.77
professor	doctor	6.62
stock	phone	1.62
stock	CD	1.31
stock	jaguar	0.92

Closet words to “Sweden” (cosine)

Word	Cosine distance

norway	0.760124
denmark	0.715460
finland	0.620022
switzerland	0.588132
belgium	0.585835
netherlands	0.574631
iceland	0.562368
estonia	0.547621
slovenia	0.531408

Word senses and word sense ambiguity

- Most words have lots of meanings!
 - Especially common words
 - Especially words that have existed for a long time
- Does one vector capture all these meanings or do we have a mess?
- Example: pike


Word senses and word sense ambiguity

pike

- A sharp point or staff
- A type of elongated fish
- A railroad line or system
- A type of road
- The future (coming down the pike)
- A type of body position (as in diving)
- To kill or pierce with a pike
- To make one's way (pike along)
- In Australian English, pike means to pull out from doing something: I reckon he could have climbed that cliff, but he piked!

Global context and multiple word prototypes [Huang et al. 2012]

- Idea: Cluster word windows around words, retrain with each word assigned to multiple different clusters bank1, bank2, etc



Linear algebraic structure of word senses

with applications to polysemy

[Arora et al. TACL 2018]

- Different senses of a word reside in a linear superposition (weighted sum) in standard word embeddings like word2vec
- $v_{\text{pike}} = a_1 v_{\text{pike1}} + a_2 v_{\text{pike2}} + a_3 v_{\text{pike3}}$
where

$$a_1 = \frac{f_1}{f_1 + f_2 + f_3}$$

- Surprising result:
 - Because of ideas from sparse coding, you can actually separate out the senses (providing they are relatively common)

Linear algebraic structure of word senses with applications to polysemy [Arora et al. TACL 2018]

tie				
trousers	season	scoreline	wires	operatic
blouse	teams	goalless	cables	soprano
waistcoat	winning	equaliser	wiring	mezzo
skirt	league	clinching	electrical	contralto
sleeved	finished	scoreless	wire	baritone
pants	championship	replay	cable	coloratura

Using word embeddings

- Word similarity: cosine
- Word clustering
- Finding similar words
- Short-document similarity
- Word analogies: $w_{\text{king}} - w_{\text{man}} + w_{\text{woman}} \approx w_{\text{queen}}$
- Odd-one out
 - Find the one that does not belong to a given list of words: computing the similarity between each word to the average word vector

Extrinsic word vector evaluation

- Extrinsic evaluation of word vectors: All subsequent tasks in this class
- One example where good word vectors should help directly: named entity recognition to identify a person, organization or location
- Next: How to use word vectors in neural net models!

Readings

1. CH 6 SLP; CH6-8, CH10-11 NNLP
2. **word2vec**: Mikolov et al. [Efficient Estimation of Word Representations in Vector Space](#). 2013
3. **Negative sampling**: Mikolov et al. [Distributed Representations of Words and Phrases and their Compositionality](#). 2013
4. **GloVe**: Pennington et al. [GloVe: Global Vectors for Word Representation](#). 2014
5. **Intrinsic and extrinsic evaluations**:
https://cs224d.stanford.edu/lecture_notes/notes2.pdf